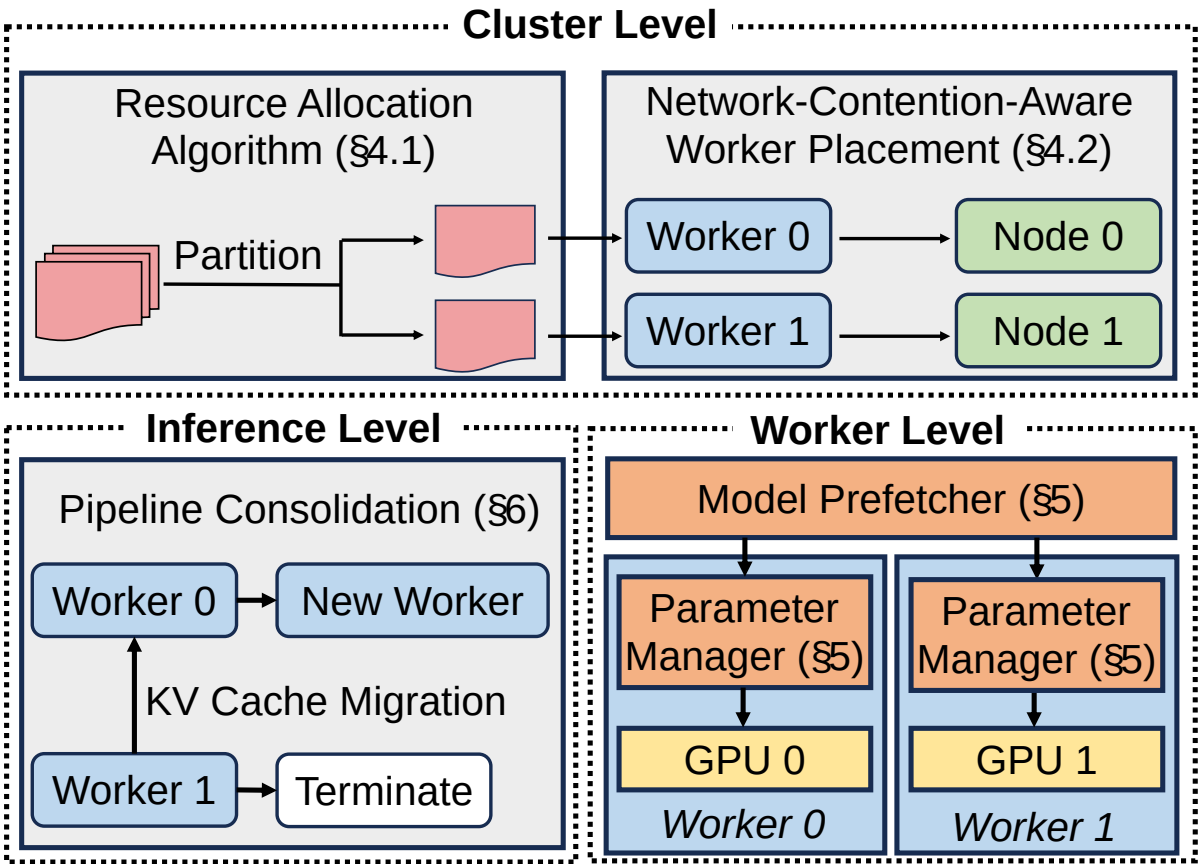
Cluster Level
Resource Allocation Algorithm (§4.1)
Partition
Network-Contention-Aware Worker Placement (§4.2)
Worker 0
Node 0
Worker 1
Node 1
Inference Level
Pipeline Consolidation (§6)
Worker 0
New Worker
KV Cache Migration
Worker 1
Terminate
Worker Level
Model Prefetcher (§5)
Parameter Manager (§5)
Parameter Manager (§5)
GPU 0
GPU 1
Worker 0
Worker 1